

Agent Skills

Skills are an open standard for extending agent capabilities. A skill is a folder containing a `SKILL.md` file with instructions that the agent can follow when working on specific tasks.

What are skills?

Skills are reusable packages of knowledge that extend what the agent can do. Each skill contains:

- **Instructions** for how to approach a specific type of task
- **Best practices** and conventions to follow
- **Optional scripts and resources** the agent can use

When you start a conversation, the agent sees a list of available skills with their names and descriptions. If a skill looks relevant to your task, the agent reads the full instructions and follows them.

Where skills live

Antigravity supports two types of skills:

Location	Scope
<code><workspace-root>/ .agents/skills/<skill-folder>/</code>	Workspace-specific
<code>~/.gemini/antigravity/skills/<skill-folder>/</code>	Global (all workspaces)

Workspace skills are great for project-specific workflows, like your team's deployment process or testing conventions.

Global skills work across all your projects. Use these for personal utilities or general-purpose tools you want everywhere.

Note: Antigravity now defaults to `.agents/skills`, but still maintains backward support for `.agent/skills`.

Creating a skill

To create a skill:

■ Agent > Skills

```
.agents/skills/  
└── my-skill/  
    └── SKILL.md
```

Every skill needs a `SKILL.md` file with YAML frontmatter at the top:

```
---  
name: my-skill  
description: Helps with a specific task. Use when you  
            need to do X or Y.  
---  
  
# My Skill  
  
Detailed instructions for the agent go here.  
  
## When to use this skill  
  
- Use this when...  
- This is helpful for...  
  
## How to use it  
  
Step-by-step guidance, conventions, and patterns the  
agent should follow.
```

Frontmatter fields

Field	Required	Description
<code>name</code>	No	A unique identifier for the skill (lowercase, hyphens for spaces). Defaults to the folder name if not provided.
<code>description</code>	Yes	A clear description of what the skill does and when to use it. This is what the agent sees when deciding whether to apply the skill.

■ Agent > Skills

While `SKILL.md` is the only required file, you can include additional resources:

```
.agents/skills/my-skill/
├── SKILL.md          # Main instructions (required)
├── scripts/          # Helper scripts (optional)
├── examples/         # Reference implementations
                        (optional)
└── resources/        # Templates and other assets
                        (optional)
```

The agent can read these files when following your skill's instructions.

How the agent uses skills

Skills follow a **progressive disclosure** pattern:

1. **Discovery:** When a conversation starts, the agent sees a list of available skills with their names and descriptions
2. **Activation:** If a skill looks relevant to your task, the agent reads the full `SKILL.md` content
3. **Execution:** The agent follows the skill's instructions while working on your task

You don't need to explicitly tell the agent to use a skill—it decides based on context. However, you can mention a skill by name if you want to ensure it's used.

Best practices

Keep skills focused

Each skill should do one thing well. Instead of a "do everything" skill, create separate skills for distinct tasks.

Write clear descriptions

The description is how the agent decides whether to use your skill. Make it specific about what the skill does and when it's useful.

Use scripts as black boxes

■ Agent > Skills

Include decision trees

For complex skills, add a section that helps the agent choose the right approach based on the situation.

Example: A code review skill

Here's a simple skill that helps the agent review code:

```
---
name: code-review
description: Reviews code changes for bugs, style
issues, and best practices. Use when reviewing PRs or
checking code quality.
---

# Code Review Skill

When reviewing code, follow these steps:

## Review checklist

1. **Correctness**: Does the code do what it's supposed
to?
2. **Edge cases**: Are error conditions handled?
3. **Style**: Does it follow project conventions?
4. **Performance**: Are there obvious inefficiencies?

## How to provide feedback

- Be specific about what needs to change
- Explain why, not just what
- Suggest alternatives when possible
```

< Rules / Workflows

Task Groups >